

# **IMPLEMENTING AND TRAINING NEURAL NETWORKS (CONT.2)**

**STFANO MELACCI  
UNIVERSITY OF SIENA**

# PYTORCH, TENSORFLOW, ...

- There exists several software libraries that can be used to easily implement Neural Networks and train them using gradient-based optimization (frequently - non always - with Python support)
  - PyTorch <https://pytorch.org/>
  - TensorFlow <https://www.tensorflow.org/>
  - scikit-learn <https://scikit-learn.org/stable/>
  - Apache MXNet <https://mxnet.apache.org/>
  - MATLAB (Deep Learning Toolbox) <https://www.mathworks.com/discovery/neural-network.html>
  - ...
- In turn, they usually exploit other lower level libraries for linear algebra or multi-threading
  - BLAS (OpenBLAS), LAPACK, Intel MKL, ATLAS, ...
  - Notice that this is also the case of NumPy

# PYTORCH, TENSORFLOW, ...

- We already discussed some NumPy-based code
  - Using NumPy we have access to common linear algebra operation and to tensor manipulation tools
- Some common questions:
  - How are these software libraries related to NumPy?
  - What do they offer that NumPy doesn't offer?
  - Which one should I use?
  - Are they so different one from each other?
  - ...

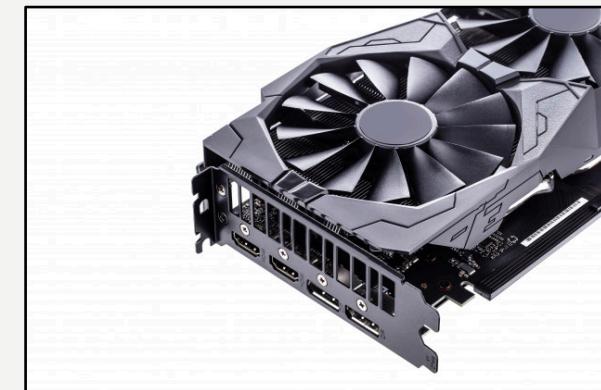


# PYTORCH, TENSORFLOW VS. NUMPY

- Let's consider PyTorch or TensorFlow as running examples
  - We will mostly consider PyTorch at a later stage
- These libraries are not so different from NumPy
  - They define their own Tensor format...
  - ...that is either fully compatible with NumPy tensors or that can be easily converted to NumPy
    - They keep a strong link with NumPy data format, providing widely used conversion tools
  - They implement fast linear algebra facilities and tensor manipulation functions
    - That are pretty similar to those of NumPy (sometimes we have very-very similar function names)
    - Parameters might have different names; with a little bit of tolerance it is easy to pass from NumPy to these libraries and vice-versa

# PYTORCH, TENSORFLOW VS. NUMPY

- However, they have some features that NumPy doesn't have!
  - The tensor data can be moved to the **GPU** memory and operations can be performed on GPUs as well
  - This is obtained with a few lines of code that specify that we are going to use a different device (we will see more details about it)
  - GPU = Fast computation on “large” scale data
    - Up to what the GPU memory can handle
  - How does one usually program GPUs?
    - NVIDIA cards and **CUDA** vs. rest of the world



# PYTORCH, TENSORFLOW VS. NUMPY

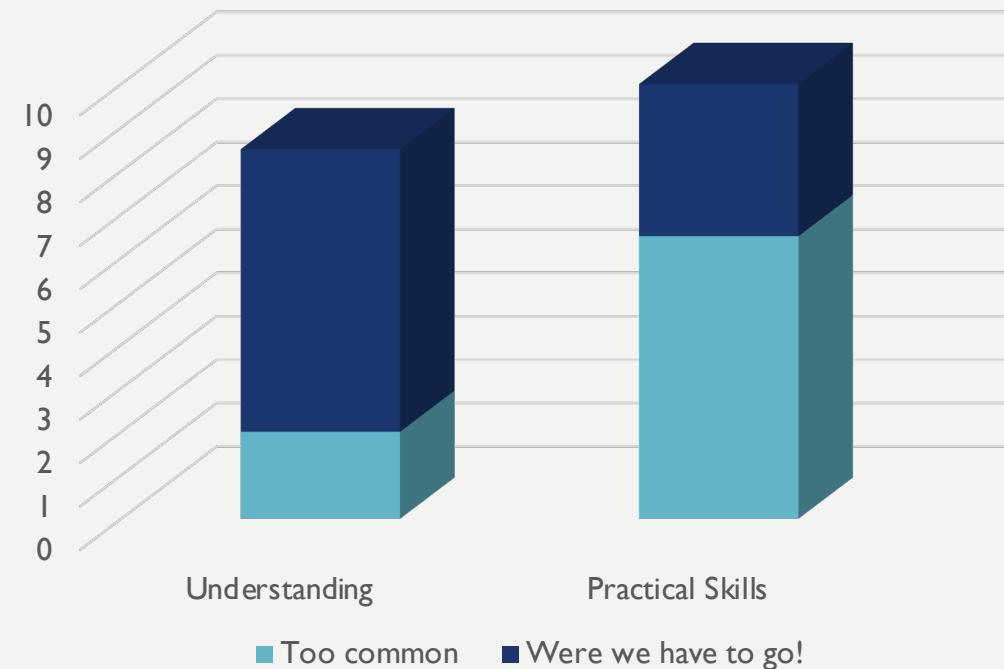
- However, they have some features that NumPy doesn't have! (cont.)
  - They perform automatic differentiation (**AutoGrad**)
    - No need to implement the chain rule 😊
    - Getting the gradients of a loss function w.r.t. to some tensors is trivial
  - How can they do that?
    - They build a **computational graph** on the fly, without showing it to the user/programmer
    - When the programmer asks for a gradient computation, they apply the chain rule traversing the computational graph in the opposite direction (as we manually did in the case of NumPy)
    - **Each of their lower level operations and functions has the derivatives w.r.t. its arguments hard-coded**
    - They keep track of the computed gradients, and, sometimes, they perform some optimization for faster computations (it is not always the case – the computational graph does introduce a small overhead)
    - *Side note: we have to talk about TensorFlow 1.x vs 2.x...*

# PYTORCH, TENSORFLOW VS. NUMPY

- However, they have some features that NumPy doesn't have! (cont. cont.)
  - They include higher level structures (classes) that can be used to easily create neural networks
  - They provide an easy-to-use interface to train a neural model, to evaluate it, hiding the math that is behind the neural-network-related computations
  - **Keras**, “*an API designed for human beings, not machines*”, has been fully integrated into *TensorFlow 2.x*  
<https://keras.io/>
  - *PyTorch* has its own interface
  - This speeds up the development of those software solutions that exploits neural network
  - Some details might get lost (major warning for newbies!)

# DEEP UNDERSTANDING VS. PRACTICAL SKILLS

- A very old story...
- Short version:
  - Gain understanding to improve your practical skills as well
    - New ideas, more energy!
  - Avoid just cloning code of other people, running it, and declaring that you are done!
    - Sometimes, too high-level software interfaces lead to this (bad) ending...



# **CODE BREAK: SLOWLY MOVING TO PYTORCH AND TENSORFLOW**

- Let's have a look at the code!
  - PyTorch-based implementation of
    - Feed-forward neural networks
    - Gradient-based weight updates
  - TensorFlow ( $\geq 2.0$ ) implementation of
    - Feed-forward neural networks
    - Gradient-based weight updates

# UNLEASH THE POWER OF GPUS!

- Small data does not allow us to see the real benefits of using GPUs to speedup computations
- They are designed to strongly parallelize larger-scale operations
  - Thousands of **parallel** threads!
  - Watch out for the way data are stored in **memory**
  - **Time** to go back and froth to/from the main (CPU) memory of the machine



*Two CPU cores vs.  
several GPU threads*



# CODE BREAK: CPU VS. GPU

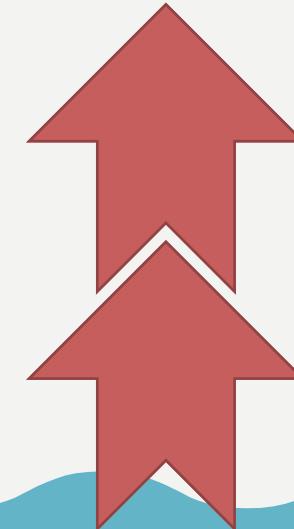
- Let's have a look at the code!
  - Matrix-by-matrix multiplication
  - CPU vs. GPU
    - NumPy, PyTorch, TensorFlow ( $\geq 2.0$ )

# NEURAL NETWORKS: PYTORCH

- Basic operations on `torch.Tensor` objects
  - `torch.op(...)`
    - <https://pytorch.org/docs/stable/torch.html>
  - `torch.nn.functional.op(...)` is a neural net-related operation
    - <https://pytorch.org/docs/stable/nn.functional.html>
- Operations as members of the `torch.Tensor` class (example: let's say we have a tensor X)
  - Usually shortcuts to basic operations
  - `X.op(...)` is sometimes equivalent to `torch.op(X, ...)`
  - `X.op_(...)` is an **in-place** operation
    - <https://pytorch.org/docs/stable/tensors.html>

# NEURAL NETWORKS: PYTORCH

- Using these operations, we can implement neural networks and all we need to do with them
  - Training, inference, ...
  - That's what we saw in the previously discussed examples!
- However, there exists **higher-level classes** that are **fully-based on these operations**, and that help in
  - *hiding the math (that's good and bad)*
  - *plugging operations together*
  - *creating more easily manageable computational blocks*
  - *handling different network layouts*
  - ...



If you look at their source code, you will see that they simply call the operations we described in the previous slide!

# NEURAL NETWORKS: PYTORCH

- Most common (base) class: `torch.nn.Module`
  - Example: a network layer is a subclass of a `torch.nn.Module`
  - Example: an MLP neural network is a set of layers, so it is a `torch.nn.Module`
- Basically, a `torch.nn.Module` collects some tensors that are **learnable variables**, and it implements a **forward** method that passes the signal through the module
  - Variables are instances of `torch.nn.parameter.Parameter`
  - They are sub-classes of `torch.Tensor`, so basically a “more decorated” `torch.Tensor`
- Modules can be collected into **containers**, that implement network layouts and pass the signal through all the collected modules (actually this is not much stressed in PyTorch as it was in Torch)
  - Containers are `torch.nn.Module` as well

# NEURAL NETWORKS: PYTORCH

- If your net is implemented as a (combination of) `torch.nn.Module`(s), then it will be easy to
  - Get a quick reference to all the network learnable parameters
  - Forward the signal through the whole network: a `torch.nn.Module` is **callable**, and, when called, it invokes the **`forward`** method of the module
  - Save and load the network
  - Use and **optimizer** to update the network parameters
    - *Stochastic Gradient Descent (SGD), Adam, ...*
  - ...



# CODE BREAK: HIDING THE MATH (PYTORCH)

- Let's have a look at the code!
  - Higher-level PyTorch-based implementation of
    - Feed-forward neural networks
    - No math involved at all (in the code)
  - (A) Using custom 'Net' class
  - (B) Using a sequential container

# NEURAL NETWORKS: TENSORFLOW

- What about TensorFlow?
  - It follows a similar organization, very similar from the lower-level perspective
  - Basic operations on tensors (as in PyTorch)
  - Higher-level modules (as `torch.nn.Module`) that are based on **Keras**
  - From this point of view, it is extremely similar to what we saw about PyTorch
  - *One of the main differences is that Keras allows us to exploit an even higher-level abstraction of the network, optimizer, accuracy measures etc... (we will not cover this at all – the focus of this course is on PyTorch)*



# CODE BREAK: HIDING THE MATH (TENSORFLOW)

- Let's have a look at the code!
  - Higher-level TensorFlow ( $\geq 2.0$ )-based implementation of
    - Feed-forward neural networks
    - No math involved at all (in the code)
  - (A) Using custom 'Net' class
  - (B) Using a sequential container

*Notice: this code could have been written using the Keras facilities, making it shorter and more compact – I kept a more verbose organization to match the PyTorch version*